

# Pythia: Compiler-Guided Defense Against Non-Control Data Attacks

Sharjeel Khan   Bodhi Chatterjee   Santosh Pande



# Software Memory Vulnerabilities = Non-Control Data Attacks




Software Memory Vulnerabilities



Privilege Escalation
Information Leakage
Control-Flow Manipulation
Unwanted Code Imports
Arbitrary Code Execution

Implications of Data Attacks




Flipping Branch Predicate by String Overflow

- ▶ Applications written in *memory unsafe languages* (C/C++) are vulnerable to data-attacks
- ▶ **Control-flow Bending** involves diverting a program branch's target into alternative paths
- ▶ Challenging to defend against because of *prevalence of program branches & program pointers*

# Conditional Branches & Program Pointers = Scalability & Analyzability

- ❑ Prevalence of conditional program branches impacts the *scalability* of this problem
  - ❑ Over **55%** of terminator IR instructions in SPEC 2017 are conditionals (~700,000)
  - ❑ Over **56%** of Nginx are conditionals (~16,000)
- ❑ Presence of program pointers affecting branches leads to *un-analyzability*
  - ❑ Out of 900,000 pointer instructions in SPEC 2017, more than 200,000 affect conditionals branches

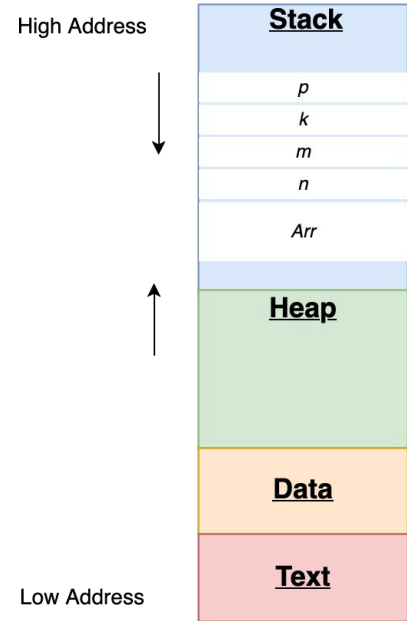


# Pointer Positioning Attack: Data Pointer Corruption

```
int *p;  
int k, m, n;  
int Arr [100];  
p = Arr;  
scanf ("% d ", &k);  
...  
m = n - 1;  
*p = n + 1;  
if ( m > n ) {  
    // privileged execution  
}
```

$p$  stores the base address of  $Arr$

$p$  points to  $m$  by overwriting pointer value and this alias sets  $m = n + 1$



Pointer **p** can be corrupted by **k** due to the overflow (memory vulnerability)

# Real-World Control-Flow Bending: Privilege Escalation

```
void Access ( char pwd [20]) {  
    char str [ SIZE ], user [ SIZE ];  
    char someinput [ DIFF_SIZE ];  
    ...  
    verify_user ( user , pwd );  
    if ( strcmp ( user , " admin ",5) ) {  
        // super user code  
        ...  
    } else {  
        // normal user code  
        ...  
    }  
  
    strcpy ( str , someinput );  
  
    if ( strcmp ( user , " admin ",5) ) {  
        // super user code  
        ...  
    } else {  
        // normal user code  
        ...  
    }  
}
```

← Input string declaration

← First string comparison after verifying username with password

← Copying str into *someinput* could lead to an overflow

← Second string comparison could be manipulated to be True, via buffer-overflow causing super user privilege

High Address

**Stack**

str

user

someinput

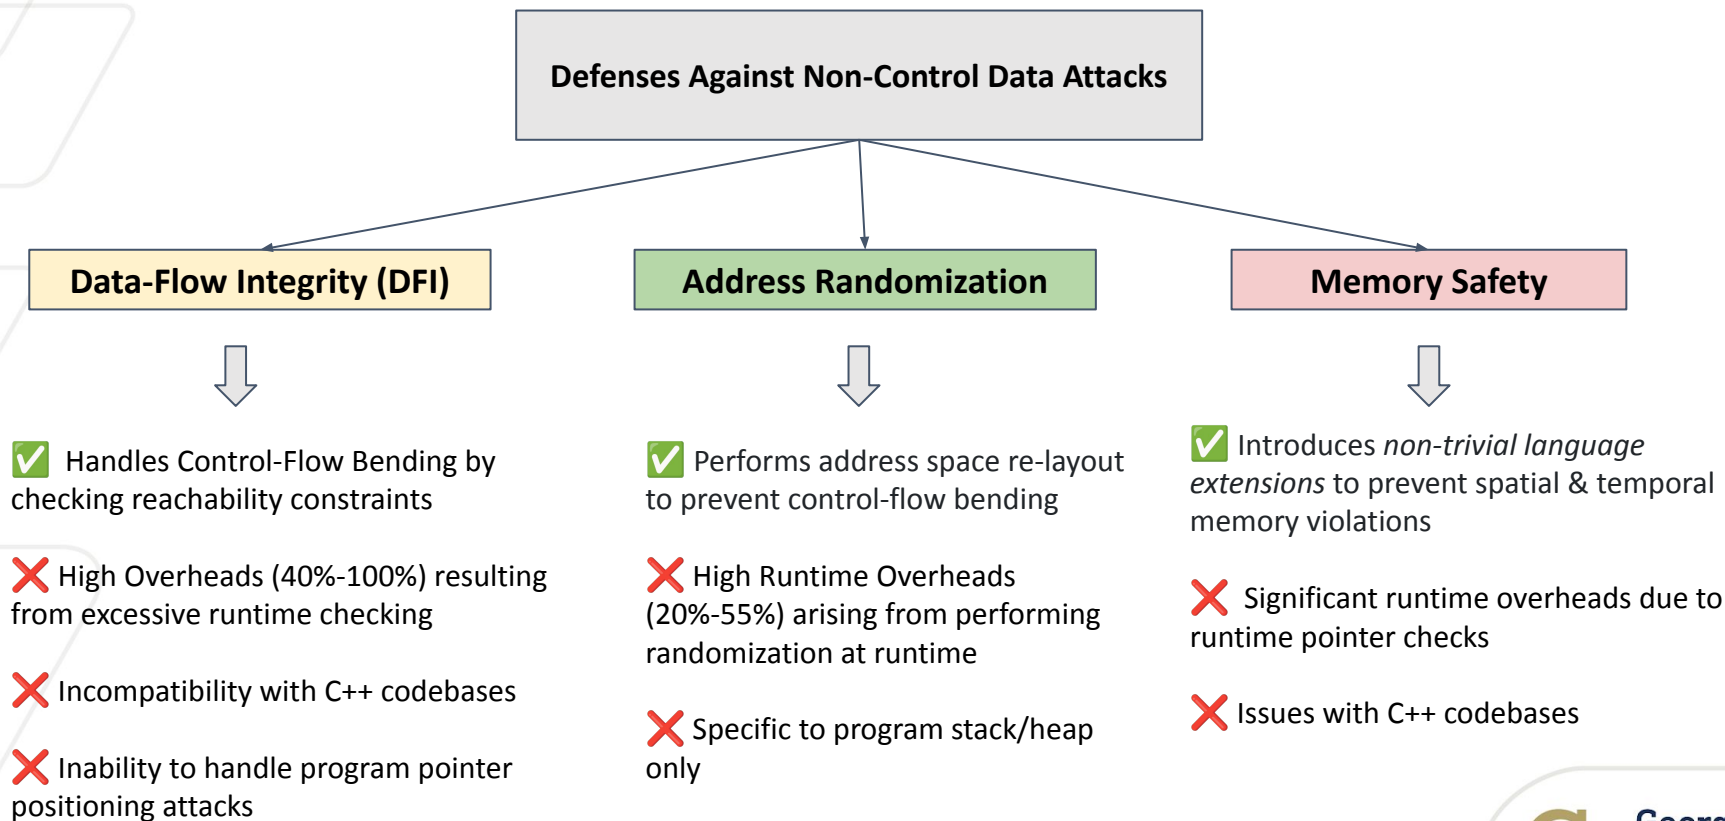
**Heap**

**Data**

**Text**


Low Address

# Existing State-of-the-Arts: High Overheads, Generalizability & Compatibility Issues

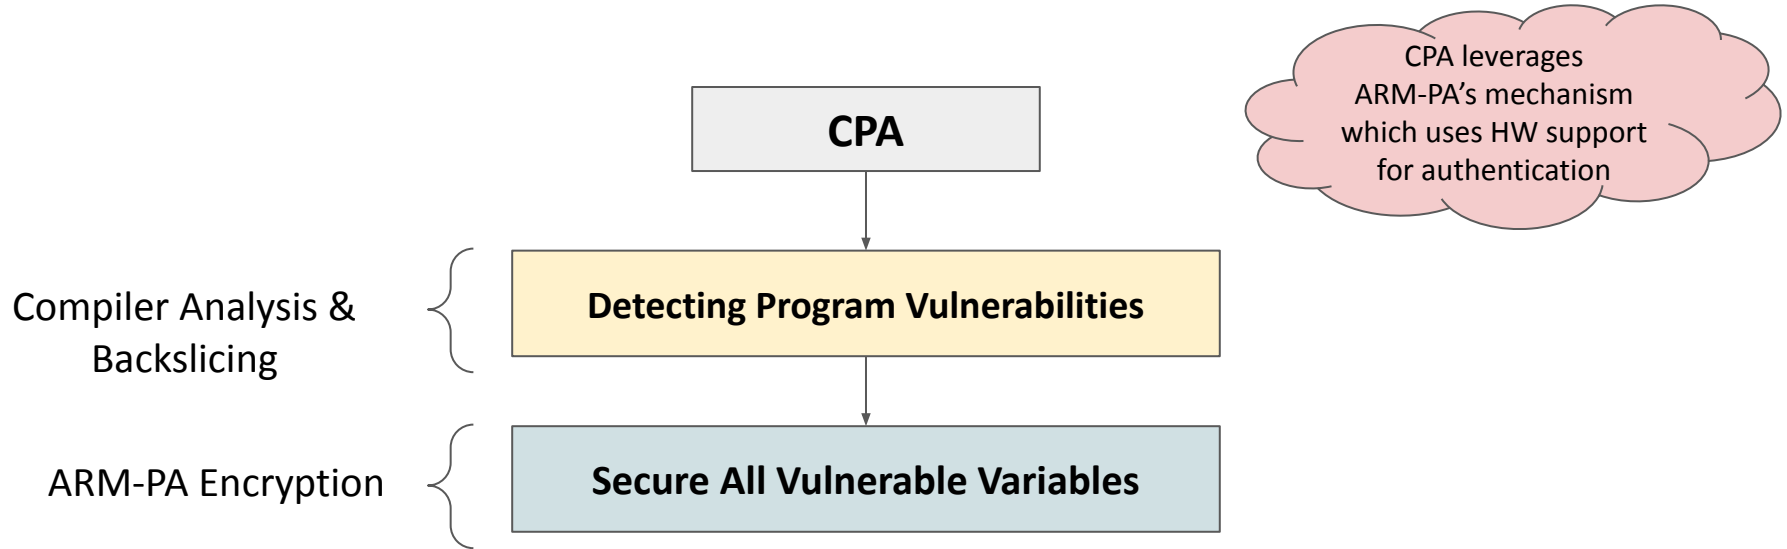


# ARM Pointer Authentication (PA)

- ❑ *Specialized hardware mechanism* that ensures the integrity of data and code pointers associated with the program
- ❑ Not all bits are required to define the address space - top bits are used as **Pointer Authentication Code (PAC)**
- ❑ In 64-bit architectures, the pointer address space is less than 40-bits so it leaves PAC with 24 bits (1 in 16 million chance)
- ❑ Each encryption (PACD\*) and authentication (AUTD\*) instruction takes around ~2ns



# Complete Pointer Authentication (CPA) for Branches




This is a baseline scheme where all possible program vulnerabilities are secured with ARM-PA to prevent overflows



# Detecting Program Vulnerabilities: Variable “Backslicing”

- ❑ For a given branch, traverse the variables’ *Use-Def* chains to detect all variables influencing that branch
- ❑ Program pointers and their possible aliases are also secured
- ❑ To secure variables:
  - ❑ **Authenticate** (decrypt) pointers variables before **loads**
  - ❑ **Sign** (encrypt) pointers in the set before **stores**

```
int a = initVar(10),
    b = initVar(20);
c = initVar(30);
d = initVar(40);
do {
    /*loop body*/
    e += a * b;
    f += c / d;
    if (e == f)
        break;
    g = a + b;
} while (g);
```



# Example: Real-World Control-Flow Bending with CPA

```
void Access ( char pwd [20]) {  
    char str [ SIZE ], user [ SIZE ];  
    char someinput [ DIFF_SIZE ];  
    llvm.pythia.sign(user);  
  
    verify_user ( user , pwd );  
    llvm.pythia.auth(user);  
    if ( strcmp ( user , " admin ",5) ) {  
        // super user code  
        ...  
    } else {  
        // normal user code  
        ...  
    }  
  
    strcpy ( str , someinput );  
    llvm.pythia.auth(user);  
    if ( strcmp ( user , " admin ",5) ) {  
        // super user code  
        ...  
    } else {  
        // normal user code  
        ...  
    }  
}
```

Encrypt pointer to user

Authenticate pointer to user

Branch Backslice Variables: user

Authenticate pointer to user

Branch Backslice Variables: user

High Address

**Stack**

str

user

someinput

**Heap**

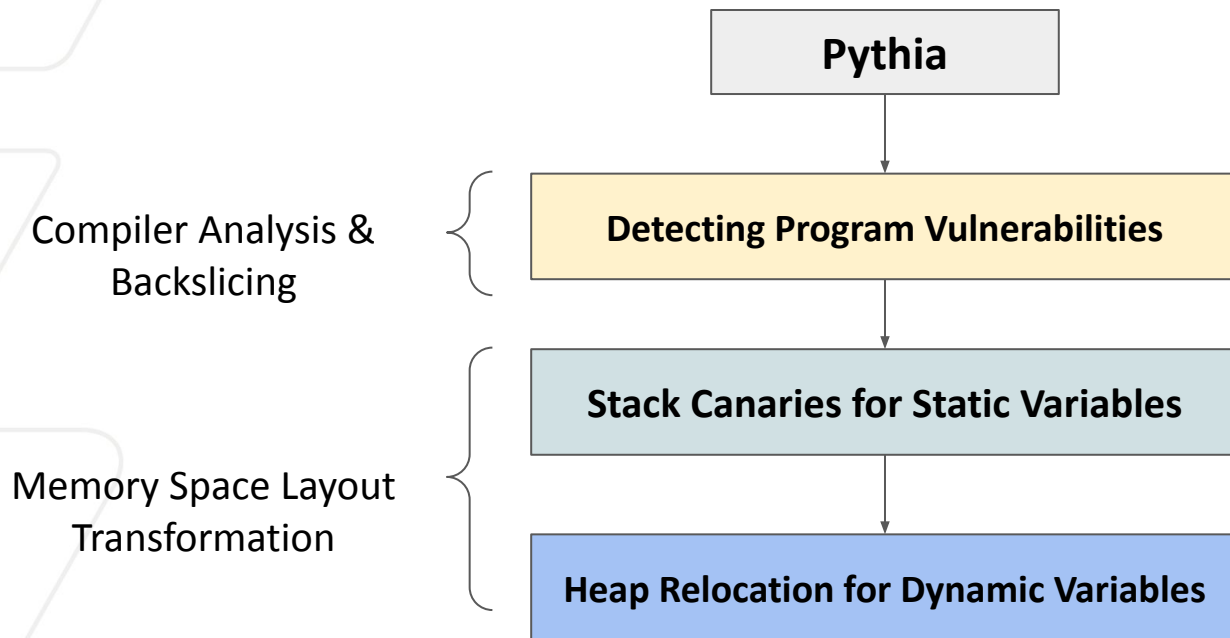
**Data**

**Text**

Low Address

*Our experiments have shown that following this naive scheme incurs upto **48%** runtime overheads, and bloats the binary size by **22%***

# Pythia: Performance-Aware Defense Approach for Control-Flow Bending

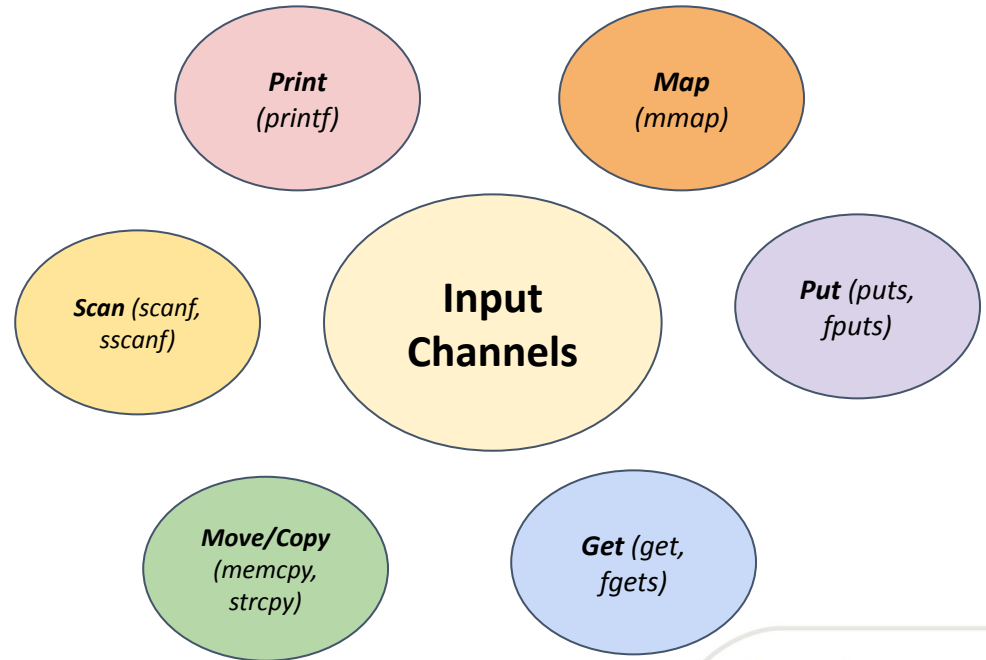


# Input Channels: Source for Attackers to Trigger Overflows

- Most overflow-based attacks are triggered via program's **input-channels**


- Attackers manipulate the arguments of input-channel functions to trigger buffer overflows

- Backslicing input-channel function variables (rather than branch variables), refines the set of vulnerable variables



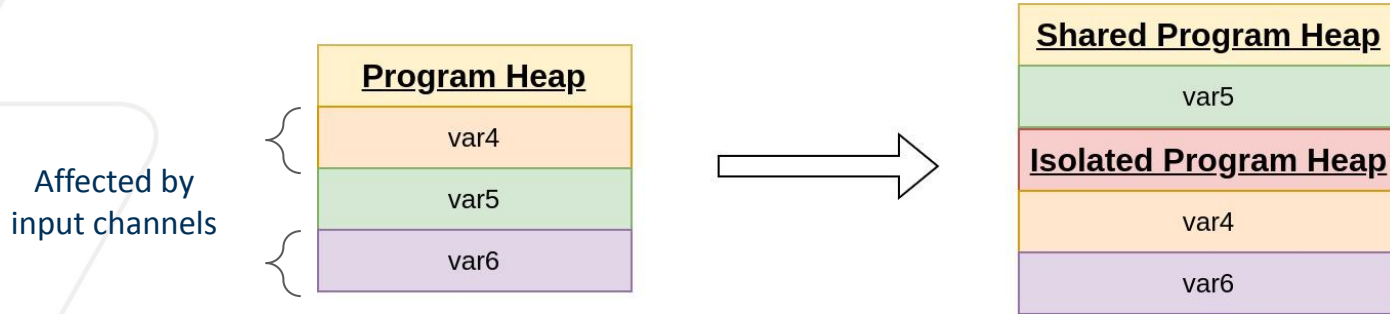
# Memory Layout Transformation: Stack Canaries

- ❑ Stack variables used in ICs are moved to the top of function frame
- ❑ Canaries are added in between each vulnerable stack variables
- ❑ Before IC for a stack variable, the canary gets re-encrypted in case of a long call-chains
- ❑ After IC for stack variable, the canary gets authenticated to check for overflows



# Memory Layout Transformation: Heap Relocation

- ❑ Vulnerable heap variable are relocated to an *isolated part of the heap* memory
- ❑ This is achieved via Pythia's custom memory allocation (e.g. malloc -> smalloc)
- ❑ Pythia's custom memory allocation reserves a big chunk of heap in the beginning and allocates based on heap allocation function calls
- ❑ These vulnerable heap variables and its uses are encrypted/authenticated with ARM PA on every store/load respectively



# Real-World Example: Control-Flow Bending with Pythia

```
void Access ( char pwd [20]) {  
    int canary = lib_random();  
    char someinput [ DIFF_SIZE ];  
    char str [ SIZE ], user [ SIZE ];  
    llvm.pythia.sign(canary);  
  
    verify_user ( user , pwd ) ;  
    if ( strcmp ( user , " admin " ,5) ) {  
        // super user code  
        ...  
    } else {  
        // normal user code  
        ...  
    }  
    llvm.pythia.auth(canary);  
    canary = lib_random();  
    llvm.pythia.sign(canary);  
    strcpy ( str , someinput ) ;  
    llvm.pythia.auth(canary);  
    ...  
}
```

← Canary got created

← someinput got moved to top

← Encrypt canary value

← Authenticate canary then re-randomize it

← Encrypt canary value

← Input Channel with someinput as target

← Authenticate canary so there was no overflow

High Address



Low Address



# Evaluation

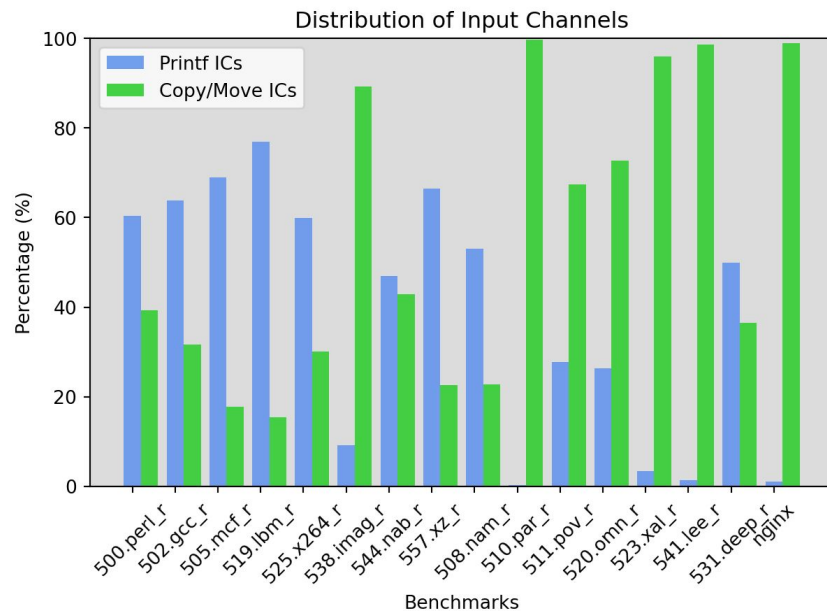
- ❑ How effective is the conservative scheme in defending against non-control data attacks, and what are its runtime overheads?
- ❑ How secure is Pythia's performance-aware approach involving stack canaries and heap sectioning approach against non-control data attacks?
- ❑ Does it manage to reduce the runtime overheads and ARM-PA instructions compared to the conservative scheme?
- ❑ How does Pythia compare to DFI in terms of securing vulnerable branches in applications that can be manipulated through the input channels?



# Summary of Results

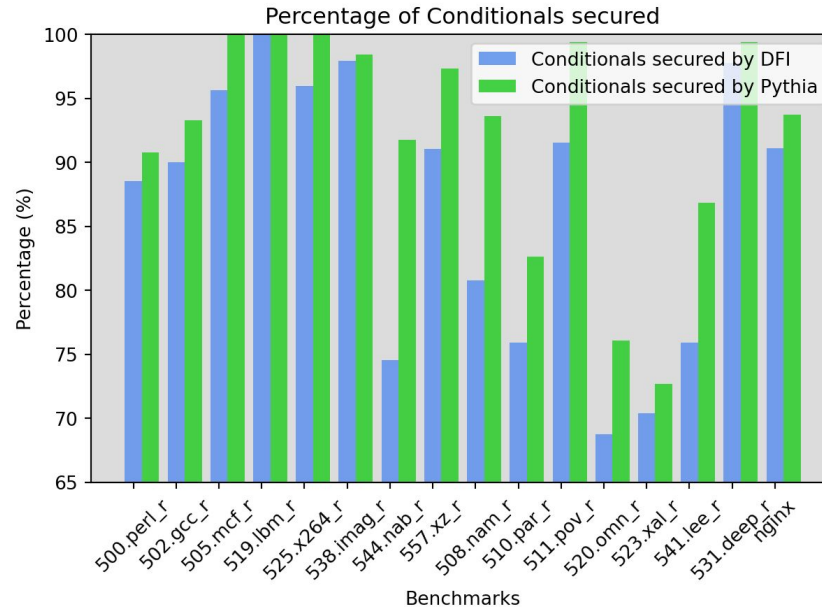
- ❑ Performance Results:
  - ❑ Average Overhead: **47.88%** (CPA) vs **13.07%** (Pythia)
  - ❑ Average Binary Increase: **21.56%** (CPA) vs **10.37%** (Pythia)
- ❑ Security Results:
  - ❑ 25326 ICs with 31.5% prints and 65.9% move/copys
  - ❑ Pythia decreased PAs by 4.25x
  - ❑ Branch protection: 92.2% (Pythia) vs 86.6% (DFI)
  - ❑ 100% Branch Protection: 3 (Pythia) vs 1 (DFI)
  - ❑ Works on well-known real-life attacks such as ProFTPD

# Security Results (Input Channels)



~25000 ICs with 31.5% prints and 65.9% move/copys


# Security Results (Branch Protection)



Branch protection: 92.2% (Pythia) vs 86.6% (DFI)

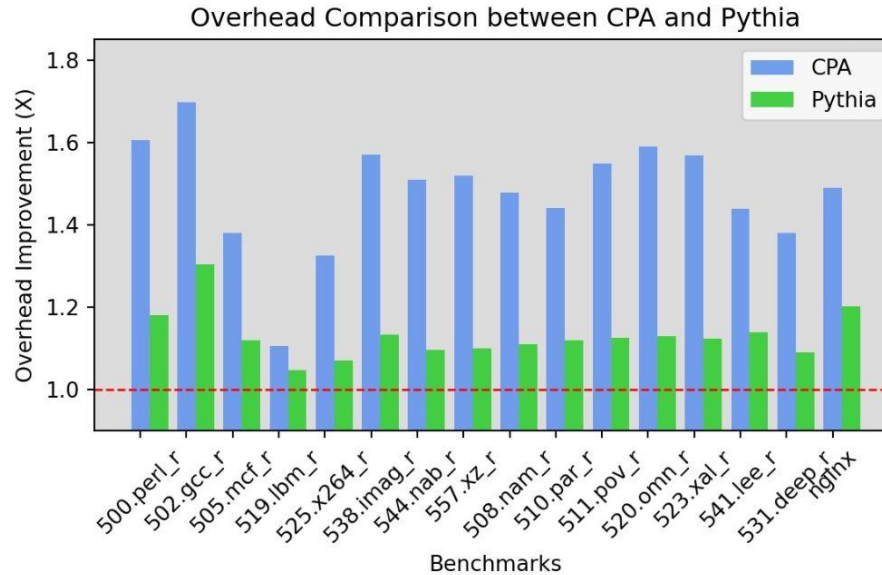
100% Branch Protection: 3 (Pythia) vs 1 (DFI)

# Performance Results (Binary Size)



Average Binary Increase: **21.56%** (CPA) vs **10.37%** (Pythia)

# Performance Results (Overheads)



Average Overhead: **47.88%** (CPA) vs **13.07%** (Pythia)  
Pythia decreased PAs by 4.25x

# Conclusion

- ❑ Pythia:
  - ❑ Compiler-guided defense framework with pointer authentication to tackle Non-Control Data Attacks
  - ❑ Utilizes hardware components already available in ARM chips
  - ❑ Effective on *pointer-intensive applications* and *C++ codes*
  - ❑ Works on well-known real-life attacks such as **ProFTPd**
- ❑ Pythia's average overhead is **13.07%** compared to Complete Pointer Authentication's average overhead of **47.88%**
- ❑ Pythia can secure **5.6% (~33000)** more branches than DFI in SPEC 2017 and Nginx and fully secure 3 applications

# Thank You

❑ **Questions?**